


Manifest — what it does, and where each service sits

Written 2026-08-11, after the first estate was deployed, verified end to end and torn down. Its purpose is to agree the shape before the upper tiers of the cascade are built.

In a paragraph


A customs broker receives commercial documents — bills of lading, invoices, packing lists, certificates of origin, customs declarations, arrival notices — as scans of scans: skewed, stamped over the numbers, in three languages, with tables that break across pages. Manifest turns each one into a **structured record where every published field points at a rectangle on a page**, checks that the documents in a shipment **agree with each other**, and proposes a **tariff classification** that a human decides. The engineering claim is not that the reading is accurate. It is that the system knows what it does not know: a field whose confidence falls below a threshold *derived from that field's declared error budget* is never published, it is sent to a person — and the queue that receives it has a declared capacity, so a system that abstains too often fails its own build rather than quietly drowning its reviewers.

1 • What happens to one document



Three things worth noticing. A document can publish *and* queue at the same time — eight fields to the record, four to a human — and getting that wrong is how the estate first shipped. And nothing reaches the record without passing the provenance gate, which checks the box against the **page**, not against the record that produced it.


2 · The boundary the whole system is built around



The four pairs are **a sample, not the list**: `src/manifest/core/` holds sixteen modules and every one of them sits in the right-hand column. The cleanest example is not in the diagram — `checkdigit.py`, which refuses an ISO 6346 container number that disagrees with **its own arithmetic**. No confidence, no threshold: either the check digit works out or it does not.

Everything on the right is pure Python in `src/manifest/core/`, importing no cloud SDK and no model library. That is what lets every claim be checked on a laptop with no AWS account — and it is why swapping Textract for a local reader is an adapter change and nothing else.

3 · The cascade — and the fact that decides its economics



Only tiers 0 and 1 report a confidence. So escalating past tier 1 is not "buying a better reading" — it is **a decision to spend a human**: the page gets read better and arrives with no score to publish it on. That is not a defect in the routing; it is what those services are, and it is why the review capacity model is costed against it rather than against an assumption.

And Greek and Dutch skip tiers 1 and 2 entirely. Textract, Bedrock Data Automation and Comprehend all document the same six input languages, and neither Greek nor Dutch is among them (docs/AWS-CONSTRAINTS.md , verified 2026-08-09). For two of this scenario's three languages, the local reader is not the cheap option — **it is the only one.**

4 · Where the AWS services actually sit


Service	Its job	Status today
S3	landing zone, records, renders, Iceberg lake	✅ running
Lambda (container)	the tier-0 reader and the provenance gate	✅ running
Step Functions	one execution per document; the escalation lives here	✅ escalation states written , never dispatched
SQS + DynamoDB	the review queue and the recorded decisions	✅ running
ECR	the reader image — 3.7 GB, two functions run <i>from</i> it	✅ running
Textract	tier 1, for the six documented languages	✅ wired , ⚠️ never called
Bedrock Data Automation	tier 2, reading order and tables	✅ wired , ⚠️ never called
Bedrock	tier 3, the only escalation for Greek and Dutch	✅ wired , ⚠️ never called
Glue · Athena · Iceberg	the record lake and its catalogue	✅ applied
OpenSearch	search over <i>published records</i> , never raw document text	✅ applied
Redshift	duty exposure by HS chapter, queue economics, cost per client	⚠️ opt-in, never applied
EMR Serverless	bulk re-extraction when the reader improves	⚠️ opt-in, never applied
SageMaker	opt-in only	⚠️ never applied

5 · What was proved, and what was not

Proved on the deployed estate, 2026-08-11, by scripts/e2e_verify.py — 13 of 13 checks. A bill of lading arrived, was read into 53 words with genuine confidences from 0.179 to 0.969 and a box for each, two fields cleared their thresholds and seven abstained, the gate verified both published fields **by cropping the page and reading it again**, the record was written keyed by document and version, and all seven abstentions reached the queue. Five edge cases were refused by name.

Not proved, and the reason is the same for all of it: no managed extraction engine has ever been called. There is no accuracy figure for the escalated fraction and there cannot be one until a page goes up a tier. No distributed job has run. Every cost figure is **modelled** from published prices, never measured.

6 · What we build next, and why it is the point



As of 2026-08-11 the escalation is written and unrun. `src/manifest/handlers/escalate.py` routes each abstaining field through `core.cascade.route`, calls the tier that field's language makes eligible, and re-thresholds only what came back with a score. It is off by default behind `enable_escalation_tiers`: with the flag down the states are not in the machine, the function does not exist, and no grant to Textract or Bedrock exists anywhere in the account.

Today the estate demonstrates tier 0 end to end and nothing above it. The routing rule is real and proved offline against the corpus; what has never happened is a page actually going up. The sentence "*the upper tiers are written and have never been called*" is honest and weak. The sentence "*a page escalated on this date, and here is what it cost*" is the one worth having — and it is a few cents of Textract away.

Redshift is the same argument for the analytics half: `contracts/` declares the marts and `scripts/check_marts.py` proves they only read columns the warehouse has, offline. Standing the workgroup up for an hour turns that from a design into an answered question — and it is cheaper than this note first claimed: an idle workgroup bills no compute at all (`docs/AWS-CONSTRAINTS.md`, read 2026-08-11). Eight RPUs for ten minutes of marts is well under a euro. What does bill continuously is a transaction nobody closed.